

Austin Varghese

Guelph, Ontario · austinfv01@gmail.com · +1-226-218-9350

Re: Operations AI Engineer — Opendoor

Dear Opendoor Hiring Team,

This letter was written by an AI agent — the same agent that generated the attached resume and submitted this application. It's available at github.com/Austin-FV/job-apply-agent. I built it because applying to the right roles felt like a problem worth automating, and because the most direct way to demonstrate that I default to AI is to let the agent speak for itself.

The stack is Python, Claude Sonnet 4.6, browser-use, and Playwright. The most interesting design problem was preventing hallucination in the resume. Every bullet is backed by a Pydantic `source_tag` that traces it to a real profile achievement — fabricated content fails schema validation before it ever hits the PDF renderer. The rest of the pipeline follows the same typed-contract logic: JD HTML gets parsed into a `JobPosting` model, the LLM produces a `ResumeContent` JSON object, Jinja2 templates it into HTML, and Playwright renders the PDF. Typed inputs and outputs at every boundary is what makes the output reliable enough to submit without a human review pass. One honest debugging note: I dropped WeasyPrint mid-build after hitting GTK install friction on Windows and switched to `page.pdf()` — Playwright was already in the stack, so the fix took minutes. The git history shows the actual path.

That build process maps directly to what this role requires. The JD calls for people who evaluate build-vs-integrate decisions, pressure-test architectures for failure modes before they become technical debt, and ship fast without sacrificing reliability. Choosing Playwright over WeasyPrint is a small example of that instinct — right tool, already present, done. The anti-hallucination schema layer is a larger one: it's the kind of observability and data integrity thinking the job description flags as non-negotiable when you're maintaining AI-powered workflows in production.

The run artifacts for this exact application — scraped JD, generated resume JSON with source tags, this cover letter, and the browser-use trace — are committed in the repo.

Best,
Austin Varghese